

## Compiler Optimization Example

Source Code looks like:

```
float A[4][4];           // A is stored in ROW major order

for (int i=0;i<4;i++) {
    A[i][j] = A[i][j]+1.0; // increment every element in jth column of A
}
```

Un-optimized assembly code looks like:

```
LD    R1,#0              // index variable i
LD    R2,j               // value of j (some constant btw 0 and 4)
LD    F1,#1.0            // floating point constant to add

LOOP: LD    R3,A          // base address of A from symbol table
      MULI  R4,R2,#16     // offset of current row
      ADD   R3,R4,R3      // R3 now has addr of first elt in row
      MUL   R4,R1,#4      // offset into row of current col
      ADD   R3,R4,R3      // R3 now has addr of A[i][j]

      LD.F  F2,(R3)
      ADD.F F2,F2,F1
      ST.F  F2,(R3)

      ADDI  R1,#1         // increment i

      SUBI  R5,R1,#4      // compare R1 and 4
      BNZ   R5,LOOP       // if R1 != 3, loop back and do another
```

The compiler can optimize away those slow multiplies in the effective address calculation

```
LD R1,#0                // R1 now holds i*4, the offset of curr elt
LD R2,j*#16             // R2 holds offset to the current row
LD F1,#1.0

LOOP: LD R3,A
      ADD R3,R3,R2
      ADD R3,R3,R1       // R3 has A + 16*j+4*i, the effective addr

      LD.F  F2,(R3)      // "meat" of the loop is unchanged
      ADD.F F2,F2,F1
      ST.F  F2,(R3)
```

```
ADD  R1,R1,#4           // increment R1 by size of an element
SUB   R5,R1,#16          // now check if R1 has 4*4 in it
BNZ   R5, LOOP
```

The important thing to note is that had the array effective address calculation not been explicit (the 5 instructions following the LOOP: label in the unoptimized version), the compiler could not have done anything with it.